

# Micro-Expression Recognition on CASME II

## Complete Results, Figures & Confusion-Matrix Walkthrough

A self-contained results document for discussion with my supervisor — every metric, every figure, and every confusion matrix explained.

### Summary of the process & technologies (in brief)

We recognise **micro-expressions** (involuntary facial movements  $< 1/5$  s) on the **CASME II** dataset. Each clip is converted into a 3-channel **motion tensor** (horizontal + vertical **optical flow** and **optical strain**), resampled to 32 frames at  $224 \times 224$ . A hybrid **PyTorch** network processes it: an optional **EVM** motion-magnification pre-step, a **three-stream 3-D CNN** spatial backbone, optional **SimAM** parameter-free attention, and an optional **SLSTT Transformer** temporal encoder, ending in a linear classifier. We run an **ablation study** — every on/off combination of these four components (12 valid configurations) — trained with **Focal loss** + class-balancing, and evaluated under two subject-disjoint protocols: **Leave-One-Subject-Out (LOSO)** and a **30 % hold-out**. The headline metric is **macro-F1**, because the 3 emotion classes (Negative / Positive / Surprise) are heavily imbalanced.

### 1. What the experiment is

Micro-expressions are brief, involuntary facial movements that leak concealed emotion. They are faint, fast, and rare, which makes recognition hard. Instead of proposing a single model, we run an **ablation study**: four components, each switched on or off, so we can measure *which parts actually help*.

| Toggle             | Component                        | What it does                                                                        |
|--------------------|----------------------------------|-------------------------------------------------------------------------------------|
| <b>EVM</b>         | Eulerian Video Magnification     | amplifies tiny, near-invisible motion before feature extraction (a data pre-step)   |
| <b>SimAM</b>       | parameter-free spatial attention | re-weights the important regions of the CNN feature maps (adds <b>0</b> parameters) |
| <b>CNN</b>         | three-stream 3-D CNN             | extracts spatial + short-range motion features — the backbone                       |
| <b>Transformer</b> | SLSTT temporal encoder           | models long-range relationships across the 32 frames                                |

Four on/off switches  $\rightarrow 2^4 = 16$  combinations; 4 are invalid (SimAM needs a CNN to attend to), leaving **12 configurations** we train and test.

**Input = motion, not pixels.** Each clip becomes a **[3, 32, 224, 224]** tensor of optical flow (u, v) + optical strain. Feeding motion — not RGB — isolates *how the skin deforms* (what a micro-expression actually is) and stops the model from recognising the person's face.

**The data.** CASME II has 255 micro-expression clips. After grouping to 3 classes and dropping the incoherent "Others" bucket, we use **156 clips from 25 subjects**:

| Class    | Clips | Built from raw emotions                                 |
|----------|-------|---------------------------------------------------------|
| Negative | 99    | disgust (63) + repression (27) + sadness (7) + fear (2) |
| Positive | 32    | happiness                                               |
| Surprise | 25    | surprise                                                |

**Why group 7 emotions into 3?** Some raw classes are unusable alone — **fear has only 2 clips, sadness only 7** — you cannot train or fairly test a 2-clip class. Negative emotions (disgust/fear/sadness/repression) also share overlapping facial signatures that even human coders confuse. Grouping by valence (Negative / Positive / Surprise) is the standard MEGC challenge protocol and keeps our numbers comparable to published work.

### 2. Why we report Macro-F1, not accuracy

The classes are very imbalanced ( $\approx 63\%$  Negative,  $20\%$  Positive,  $16\%$  Surprise). A lazy model that **always predicts "Negative"** and learns nothing scores:

| "Always predict Negative" | Accuracy     | Macro-F1 |
|---------------------------|--------------|----------|
| LOSO                      | <b>0.662</b> | 0.266    |
| Hold-out                  | <b>0.750</b> | 0.286    |

So a useless model already gets **66–75 % accuracy**. Optimising for accuracy would *reward* ignoring the rare classes. **Macro-F1** averages each class's F1 equally, so neglecting Positive/Surprise is punished — it is the honest metric here. We therefore lead with macro-F1, keep accuracy only as a reference, and report per-class precision/recall/F1 and the full confusion matrices to show *where* errors occur.

### 3. The two validation methods (brief)

Both are **subject-disjoint** — the same person is *never* in both training and testing. This prevents **identity leakage** (a model recognising the *face* instead of the *expression*, which fakes the score).

- **Hold-out (fast):** split subjects **once** —  $\sim 70\%$  train, **30 % held out for testing** (52 test clips). One run per config; used to iterate. Downside: depends on the luck of one split (higher variance).
- **LOSO — Leave-One-Subject-Out (rigorous):** hold out **one subject**, train on the rest, test on that subject, repeat, pool predictions. It measures performance on a **completely new person** — the number that matters for real use. Downside: one full run per subject  $\rightarrow$  very slow, so we ran a **pilot of 20 of the 25 subjects** (139 pooled test clips).

**Why 30 % held out (and not more)?** With only  $\sim 140$  clips, a  $10\%$  test set ( $\sim 15$  clips) is too small to measure a 3-class score and might contain zero Surprise clips.  $30\%$  ( $\sim 52$  clips) is the smallest test set that keeps all three classes present and measurable. Testing on  $100\%$  is impossible — nothing would be left to train on.

## 4. Results tables

### 4.1 LOSO (pilot, 20/25 subjects, 139 test clips) — ranked by Macro-F1

| Rank | Configuration              | EVM | SimAM | CNN | Transf. | Accuracy | Macro-F1     |
|------|----------------------------|-----|-------|-----|---------|----------|--------------|
| 1    | attention_base             | –   | ✓     | ✓   | –       | 0.643    | <b>0.379</b> |
| 2    | spatial_only               | –   | –     | ✓   | –       | 0.627    | 0.358        |
| 3    | EVM + CNN                  | ✓   | –     | ✓   | –       | 0.491    | 0.329        |
| 4    | EVM + SimAM + CNN          | ✓   | ✓     | ✓   | –       | 0.491    | 0.324        |
| 5    | full_no_attention          | ✓   | –     | ✓   | ✓       | 0.419    | 0.264        |
| 6    | EVM + Transformer          | ✓   | –     | –   | ✓       | 0.511    | 0.256        |
| 7    | CNN + Transformer          | –   | –     | ✓   | ✓       | 0.476    | 0.252        |
| 8    | temporal_only              | –   | –     | –   | ✓       | 0.600    | 0.251        |
| 9    | full (no EVM)              | –   | ✓     | ✓   | ✓       | 0.474    | 0.250        |
| 10   | <b>proposed (all 4 on)</b> | ✓   | ✓     | ✓   | ✓       | 0.413    | 0.249        |
| 11   | pure_base                  | –   | –     | –   | –       | 0.295    | 0.179        |
| 12   | motion_amp_base            | ✓   | –     | –   | –       | 0.305    | 0.166        |
| —    | <i>do-nothing baseline</i> |     |       |     |         | 0.662    | 0.266        |

### 4.2 Hold-out (30 % subjects, 52 test clips) — ranked by Macro-F1

| Rank | Configuration     | EVM | SimAM | CNN | Transf. | Accuracy | Macro-F1     |
|------|-------------------|-----|-------|-----|---------|----------|--------------|
| 1    | EVM + CNN         | ✓   | –     | ✓   | –       | 0.577    | <b>0.458</b> |
| 2    | EVM + SimAM + CNN | ✓   | ✓     | ✓   | –       | 0.558    | 0.448        |
| 3    | attention_base    | –   | ✓     | ✓   | –       | 0.481    | 0.402        |

| Rank | Configuration              | EVM | SimAM | CNN | Transf. | Accuracy | Macro-F1 |
|------|----------------------------|-----|-------|-----|---------|----------|----------|
| 4    | spatial_only               | –   | –     | ✓   | –       | 0.462    | 0.388    |
| 5    | CNN + Transformer          | –   | –     | ✓   | ✓       | 0.442    | 0.274    |
| 6    | temporal_only              | –   | –     | –   | ✓       | 0.231    | 0.247    |
| 7    | full (no EVM)              | –   | ✓     | ✓   | ✓       | 0.308    | 0.218    |
| 8    | EVM + Transformer          | ✓   | –     | –   | ✓       | 0.212    | 0.210    |
| 9    | <b>proposed (all 4 on)</b> | ✓   | ✓     | ✓   | ✓       | 0.173    | 0.183    |
| 10   | full_no_attention          | ✓   | –     | ✓   | ✓       | 0.192    | 0.176    |
| 11   | pure_base                  | –   | –     | –   | –       | 0.192    | 0.111    |
| 12   | motion_amp_base            | ✓   | –     | –   | –       | 0.096    | 0.129    |
| —    | <i>do-nothing baseline</i> |     |       |     |         | 0.750    | 0.286    |

### 4.3 Four takeaways

1. **The 3-D CNN is essential.** Top-4 configs (both protocols) all have CNN✓ and TransformerX. Every CNN-less config sinks to the bottom (EVM-only bottoms at 0.096 hold-out accuracy).
2. **The Transformer hurts.** In every matched pair, turning it on *lowers* macro-F1 (LOSO config\_5→6: –0.13; hold-out config\_13→7: –0.28). Cause: a 4-layer/8-head encoder (~10<sup>5</sup> parameters) overfits on ~140 clips.
3. **More ≠ better.** The fully-loaded proposed model (all 4 on) ranks 10/12 (LOSO) and 9/12 (hold-out) — an honest negative result.
4. **EVM is protocol-dependent.** It helps hold-out (top-2 use it) but the best LOSO model omits it — magnification amplifies motion *and* noise, and the noise does not transfer to unseen subjects.

## 5. The summary figures explained


Each figure exists for both protocols ([losa/plots/](#) and [holdout/plots/](#)).

### 5.1 Accuracy vs Macro-F1 (bar chart)


**Axes:** x = the 12 configurations; y = score 0–1. Blue = accuracy, orange = macro-F1.

**What to see:** the blue (accuracy) bar is **always taller** than the orange (macro-F1) bar. That persistent gap is the fingerprint of majority-class bias — models get many Negative clips right (high accuracy) but do poorly on the rare classes (low macro-F1). This chart is the single clearest justification for leading with macro-F1.

**LOSO:**



**Hold-out:**




## 5.2 Per-class F1 (grouped by class)

**Axes:** x = the three classes (Negative / Positive / Surprise); y = F1 0–1; one coloured bar per configuration.


**What to see:** Negative F1 is high and uniform (every model finds the majority class). Positive is middling and is where good configs separate. Surprise is the lowest and most scattered — the hardest, rarest class. The ranking is decided almost entirely on Positive and Surprise, and the fact that minority bars are not flat-zero proves the imbalance defences are working. In LOSO, note config\_4's near-zero Negative bar — EVM-only, CNN-less features fail even on the easy class.

**LOSO:**



LOSO — Per-class F1 by configuration

Hold-out:




Hold-out — Per-class F1 by configuration

### 5.3 All confusion matrices (3×4 grid)

**Axes (each small matrix):** rows = true class, columns = predicted class (N / P / S); darker = more clips.


**What to see:** a bright **diagonal** = good; a bright **vertical stripe** = the model collapsed onto one predicted class. Good configs (3, 5, 13, 16) show a diagonal (strong top-left Negative cell). Degenerate / Transformer-heavy configs (1, 8, 4, 2) show vertical banding — e.g. config\_1 predicts *no* clip as Negative (empty first column). This is the mechanism behind the F1 numbers: failures are systematic column-collapse, not random error.

LOSO:



LOSO — Confusion matrices, all 12 configs

Hold-out:




Hold-out — Confusion matrices, all 12 configs

## 5.4 Key configurations side-by-side

**What to see:** the headline configs placed adjacently. The proposed all-in model's **smeared** matrix sits next to the **clean diagonal** of the lean CNN configs — making the negative result (adding EVM + Transformer degrades a working CNN


model) visually undeniable.

## LOSO:



LOSO — Key configs confusion, side by side

## Hold-out:



Hold-out — Key configs confusion, side by side

## 6. Every confusion matrix, read one by one

Below, each configuration's confusion matrix is given as three rows — the **true** class — with the three numbers being predictions for **N**, **P**, **S**. "Reading" explains the failure/success mode. (A perfect model would have all mass on the diagonal: N→N, P→P, S→S.)

### 6.1 LOSO — numeric confusion matrices (139 clips)

| Config (E/S/C/T)        | Acc       | F1        | True N<br>→(N,P,S) | True P<br>→(N,P,S) | True S<br>→(N,P,S) | Reading                                                                                                       |
|-------------------------|-----------|-----------|--------------------|--------------------|--------------------|---------------------------------------------------------------------------------------------------------------|
| 1 pure_base (....)      | 0.29<br>5 | 0.17<br>9 | 0,33,59            | 0,22,8             | 0,4,13             | <b>Total majority collapse</b> — predicts <i>nothing</i> as Negative (empty N column). Floor baseline.        |
| 2 temporal_only (...T)  | 0.60<br>0 | 0.25<br>1 | 31,40,2<br>1       | 3,23,4             | 4,4,9              | Over-predicts Positive; high accuracy but Negative recall only 0.34 — accuracy inflated by majority guessing. |
| 3 spatial_only (..C.)   | 0.62<br>7 | 0.35<br>8 | 47,15,3<br>0       | 4,17,9             | 1,4,12             | <b>Clear diagonal.</b> Strong Negative (47), leaks 30 to Surprise. Best simple CNN.                           |
| 4 motion_amp (E...)     | 0.30<br>5 | 0.16<br>6 | 1,55,36            | 0,27,3             | 0,5,12             | <b>Worst.</b> EVM without a CNN predicts almost everything Positive/Surprise; Negative recall 0.01.           |
| 5 attention_base (.SC.) | 0.64<br>3 | 0.37<br>9 | 50,11,3<br>1       | 6,15,9             | 2,1,14             | <b>Best LOSO.</b> Strongest Negative diagonal (50), Surprise recall 0.82. Cleanest matrix.                    |
| 6 full noEVM (.SCT)     | 0.47<br>4 | 0.25<br>0 | 31,34,2<br>7       | 6,20,4             | 4,1,12             | Adding the Transformer to config_5 <b>smears</b> the Negative row (recall 0.54→0.34).                         |
| 7 full_no_attn (E.CT)   | 0.41<br>9 | 0.26<br>4 | 19,28,4<br>5       | 5,22,3             | 0,3,14             | Negative badly leaked to Surprise (45); recall 0.21. EVM+Transformer hurts.                                   |
| 8 proposed (ESCT)       | 0.41<br>3 | 0.24<br>9 | 20,34,3<br>8       | 6,22,2             | 0,5,12             | The full model <b>smears Negative</b> across Positive(34)/Surprise(38); recall 0.22.                          |
| 9 CNN+Transf (..CT)     | 0.47<br>6 | 0.25<br>2 | 30,31,3<br>1       | 7,20,3             | 4,2,11             | Negative split almost evenly across all three — classic Transformer smear.                                    |
| 12 EVM+Transf (E..T)    | 0.51<br>1 | 0.25<br>6 | 18,50,2<br>4       | 2,24,4             | 2,5,10             | Over-predicts Positive (50); Negative recall 0.20. No CNN → weak.                                             |
| 13 EVM+CNN (E.C.)       | 0.49<br>1 | 0.32<br>9 | 42,21,2<br>9       | 6,21,3             | 1,1,15             | Good diagonal; <b>best minority F1 of the EVM configs</b> ; Surprise recall 0.88.                             |
| 16 EVM+SimAM+CNN (ESC.) | 0.49<br>0 | 0.32<br>4 | 41,22,2<br>9       | 5,22,3             | 0,2,15             | Nearly identical to config_13 — SimAM adds little on top of EVM+CNN here.                                     |

### 6.2 Hold-out — numeric confusion matrices (52 clips)

| Config (E/S/C/T)        | Acc       | F1        | True N<br>→(N,P,S) | True P<br>→(N,P,S) | True S<br>→(N,P,S) | Reading                                                                        |
|-------------------------|-----------|-----------|--------------------|--------------------|--------------------|--------------------------------------------------------------------------------|
| 1 pure_base (....)      | 0.19<br>2 | 0.11<br>1 | 0,37,2             | 0,10,0             | 0,3,0              | Predicts almost everything Positive; Negative & Surprise recall 0. Degenerate. |
| 2 temporal_only (...T)  | 0.23<br>1 | 0.24<br>7 | 0,30,9             | 0,9,1              | 0,0,3              | No Negative predicted (empty N column) — collapse.                             |
| 3 spatial_only (..C.)   | 0.46<br>2 | 0.38<br>8 | 18,5,16            | 4,3,3              | 0,0,3              | Diagonal present; Surprise recall 1.0 but precision 0.14 (over-predicts S).    |
| 4 motion_amp (E...)     | 0.09<br>6 | 0.12<br>9 | 0,3,36             | 0,2,8              | 0,0,3              | <b>Worst (acc 0.096).</b> Predicts nearly everything Surprise.                 |
| 5 attention_base (.SC.) | 0.48<br>1 | 0.40<br>2 | 19,4,16            | 4,3,3              | 0,0,3              | Like config_3, slightly stronger Negative. Solid diagonal.                     |
| 6 full noEVM (.SCT)     | 0.30<br>8 | 0.21<br>8 | 11,24,4            | 3,5,2              | 0,3,0              | Transformer smears Negative into Positive(24); <b>Surprise F1 = 0.</b>         |




| Config<br>(E/S/C/T)         | Acc       | F1        | True N<br>→(N,P,S) | True P<br>→(N,P,S) | True S<br>→(N,P,S) | Reading                                                                                                  |
|-----------------------------|-----------|-----------|--------------------|--------------------|--------------------|----------------------------------------------------------------------------------------------------------|
| 7 full_no_attn<br>(E.CT)    | 0.19<br>2 | 0.17<br>6 | 0,34,5             | 0,9,1              | 0,2,1              | Negative collapse (recall 0); predicts Positive.                                                         |
| 8 <b>proposed</b><br>(ESCT) | 0.17<br>3 | 0.18<br>3 | 0,15,24            | 0,7,3              | 0,1,2              | <b>The full model predicts NO Negative at all</b> (recall 0) — worst-case majority collapse on hold-out. |
| 9 CNN+Transf<br>(..CT)      | 0.44<br>2 | 0.27<br>4 | 20,0,19            | 6,0,4              | 0,0,3              | Predicts <b>no Positive</b> (empty P column); over-predicts Surprise.                                    |
| 12<br>EVM+Transf<br>(E..T)  | 0.21<br>2 | 0.21<br>0 | 0,36,3             | 0,10,0             | 0,2,1              | Negative collapse; predicts Positive.                                                                    |
| 13 EVM+CNN<br>(E.C.)        | 0.57<br>7 | 0.45<br>8 | 24,13,2            | 5,5,0              | 0,2,1              | <b>Best hold-out.</b> Strong Negative diagonal (24), balanced across classes.                            |
| 16 EVM+SimA<br>M+CNN (ESC.) | 0.55<br>8 | 0.44<br>7 | 23,14,2            | 5,5,0              | 0,2,1              | Near-identical to config_13 — again SimAM $\approx$ neutral on top of EVM+CNN.                           |


**Cross-cutting pattern:** every matrix with an **empty column** is a collapsed model (it never predicts that class), and these are exactly the CNN-less or Transformer-heavy configs. Every matrix with a **populated diagonal** has the CNN on and the Transformer off. This is the same story as the ranking tables, now visible at the level of individual predictions.

## 7. Appendix — individual per-config confusion matrices


Full-resolution confusion matrix for each configuration (also stored at `{loso,holdout}/config_*/confusion_matrix.png`).

### 7.1 LOSO







config\_7 full\_no\_attn — 0.419 / 0.264




config\_8 proposed (all 4) — 0.413 / 0.249




config\_9 CNN+Transf — 0.476 / 0.252



config\_12 EVM+Transf — 0.511 / 0.256




config\_13 EVM+CNN — 0.491 / 0.329




config\_16 EVM+SimAM+CNN — 0.490 / 0.324


## 7.2 Hold-out




config\_1 pure\_base — 0.192 / 0.111




config\_2 temporal\_only — 0.231 / 0.247




config\_3 spatial\_only — 0.462 / 0.388




config\_4 motion\_amp — 0.096 / 0.129




config\_5 attention\_base — 0.481 / 0.402




config\_6 full noEVM — 0.308 / 0.218




config\_7 full\_no\_attn — 0.192 / 0.176




config\_8 proposed (all 4) — 0.173 / 0.183




config\_9 CNN+Transf — 0.442 / 0.274



config\_12 EVM+Transf — 0.212 / 0.210



config\_13 EVM+CNN (BEST) — 0.577 / 0.458



config\_16 EVM+SimAM+CNN — 0.558 / 0.447

## 8. Key settings (reference)

- **Loss:** Focal loss ( $\gamma = 2$ ) + label smoothing 0.05 + inverse-frequency class weights + balanced oversampling (three stacked defences against the 63/20/16 imbalance).
- **Optimiser:** AdamW, lr 1e-4, weight decay 1e-4, gradient clipping 1.0, mixed precision, CosineAnnealing schedule.
- **Batch size 2** — a hardware limit (a [3,32,224,224] clip needs ~5 GB VRAM for Transformer configs).
- **Best checkpoint selected by macro-F1** (not accuracy); seed 42 for reproducibility.

All numbers from `results_weekend/{loso,holdout}/summary.csv` and each config's `final_results.json`; class counts from `master_thesis_labels.csv`; 30 % hold-out from `gui_settings.json`.